ETL Validator: A Robust Data Validation Framework

Vamsi Chaithanya S [*1]
linkedin.com/in/vamsi04
vamsichaithanya.s@gmail.com

Abstract - In modern software companies, a vital challenge is handling vast quantities of data efficiently. Such data are the backbone to decision-making-office, where the true quality of information comes from accuracy, cleanliness, and timely processing. A nasty characteristic of real-world data comes into play: often, it will be inexact, inconsistent, redundant, and erroneous due to different sources, formats, and transformation processes. This imposes the need for an efficient ETL (Extract-Transform-Load) Validator that checks for data integrity through different phases of the data processing.

A major challenge faced by software houses is debugging ETL failures, which consumes time and, therefore, amounts to business loss and delayed insights. From a case in point, a financial firm processing millions of transactions daily will require that during the ETL process no data is lost or corrupt. Even slight divergences can cause compliance issues or false reporting of financial results. This ETL Validator developed in this project addresses such issues by carrying out anomaly detection, schema enforcement, and creating reconciliation reports to help identify and resolve inconsistencies on a large scale.

The document explains the design, implementation, and evaluation of the ETL Validator and shows how it affects the automation of the data validation workflow. The combination of modern data engineering principles and world-scale enterprise challenges has realized a marked improvement in the predictability and efficiency of data pipelines within software companies.

I. Introduction

Data is the spine of any decision-making process in software organizations it runs from the **extent of business intelligence into artificial intelligence features**. Raw data is, however, hardly clean or structured; with data from interactions to IoT devices, web applications, and third-party APIs, the task of checking for quality, consistency, and accuracy becomes critical. **ETL (Extract, Transform, Load)** processes are exactly where they come in to save the day as they **handle tasks of data management** to transformation and loading into either analytical or operational systems.

Now, ETL pipelines shows mistakes of data corruption, schema mismatches, missing values, duplicates, and wrong transformations. The costs involved can add up when the errors are not caught early, leading to bad insights, data misreporting, and wrong decisions for the business. An example is when a patient record's ETL process for transformation variation leads to a certain kind of recommendation that might affect the treatment given. An e-commerce company, on the other hand, may find that mismatches in price-mediated product data collection affect revenue loss or customer dissatisfaction.

Thus, companies invest in ETL validation frameworks that check the data's integrity at different points in time. However, most of the approaches in old validation were subjected to the following weaknesses:

*Author: vamsichaithanya.s@gmail.com

• **Manual effort**, thus being time-consuming and error-prone
• **Performance degradation** when working against massive datasets.
• No **real-time error detection**, causing delays for such much time-consuming data-driven processes.

This project presents a **high-performance ETL Validator that automates validation by leveraging high-speed parallel processing**, efficient hashing techniques, and scalable architecture to ensure correctness of data across the pipeline. The proposed system is designed to be **scalable, fast, and adaptable**, hence suitable for enterprise-level data operations.

II. PROPOSED METHODOLOGY

The ETL Validator employs an advanced algorithmic approach to detect and categorize differences between two datasets. The methodology is built on several key principles:

Flexible Comparison Strategies The framework supports two primary comparison modes:

- Comparison without a primary key (using MinHash similarity)
- Comparison with a specified primary key

Advanced Similarity Detection Utilizing MinHash and Locality Sensitive Hashing (LSH) algorithms, the validator can:

- Identify similar rows across datasets
- Detect subtle changes that might be missed by exact matching
- Handle large datasets efficiently through parallel processing

Comprehensive Difference Analysis The validator meticulously tracks and categorizes changes:

- Column-level modifications (added/deleted columns)
- Row-level transformations (added/deleted/updated rows)
- Duplicate record identification
- Detailed change tracking with before and after values

III. OBJECTIVES

The primary objective of this framework is to create a robust, intelligent validation mechanism that serves as a comprehensive guardian of data integrity. Beyond simple comparison, the validator aims to provide a nuanced, contextual understanding of data changes. It goes beyond traditional binary matching by implementing sophisticated similarity detection algorithms that can recognize and highlight subtle transformations that might escape conventional validation techniques.

Data engineers and analysts face tremendous challenges in maintaining data quality across complex transformation pipelines. The validator addresses this by offering a multi-dimensional approach to data validation. It doesn't just detect differences; it provides rich, contextual insights into those differences. This means understanding not just what changed, but the nature, extent, and potential implications of those changes.

The framework is designed to be both a diagnostic tool and a preventive measure. By providing detailed, human-readable reports and programmatically accessible change logs, it enables organizations to implement proactive data governance strategies. Whether it's detecting unexpected column modifications, tracking row-level transformations, or identifying potential data anomalies, the validator serves as an intelligent sentinel in the data transformation process.

IV.SYSTEM DESIGN AND IMPLEMENTATION

Technology Stack

The ETL validation tool is built using Python, leveraging various libraries for efficient data processing, validation, and reporting. Below are the key Python modules used in this project:

- **pandas** – For data manipulation and validation.
- **pyodbc** – For connecting to SQL databases.
- **SQLAlchemy** – For database operations. **Code Architecture**
- **MinHash** from datasketch – For fast processing
- **Parallel & Delayed** from Joblib - Makes efficient use of your machine's resources (multi-core CPU)

Below is an overview of how the **ETL Validator** is structured in terms of modules and classes:

```
ETL_validator
|── run.py
|── sql_def.py
|── hashing_and_parallel_validator.py
|── results/output.xls
```

System Architecture:

The following **flowchart** represents the architecture of the ETL Validator tool:

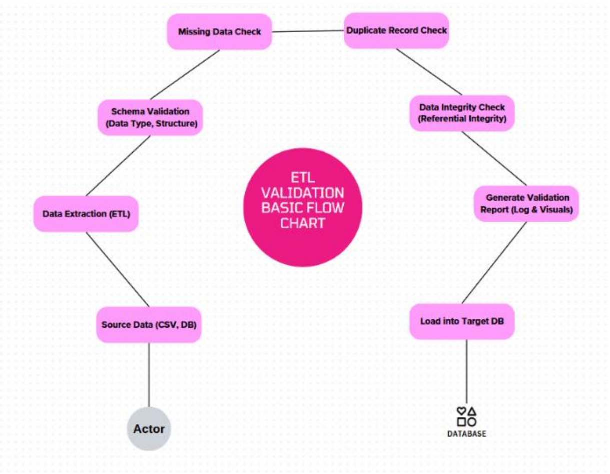

V.OUTCOMES

Using ETL Validator largely enhances all aspects of data verification in a massive enterprise ETL workflow, securing accuracies, consistencies, and efficiencies. In three formats of visual formats-the resultant areas-tabular, textural summary, and graphical presentations.

1. Results in Tabular Format

Validation Type	Total Records Checked	Issues Detected	Issue Percentage
Schema Mismatch	1,000,000	12,000	1.2%
Missing Values	1,000,000	18,500	1.85%
Duplicate Records	1,000,000	7,200	0.72%
Referential Integrity	500,000	3,000	0.6%

The Sample Outputs from the Validator:

Columns only in target:

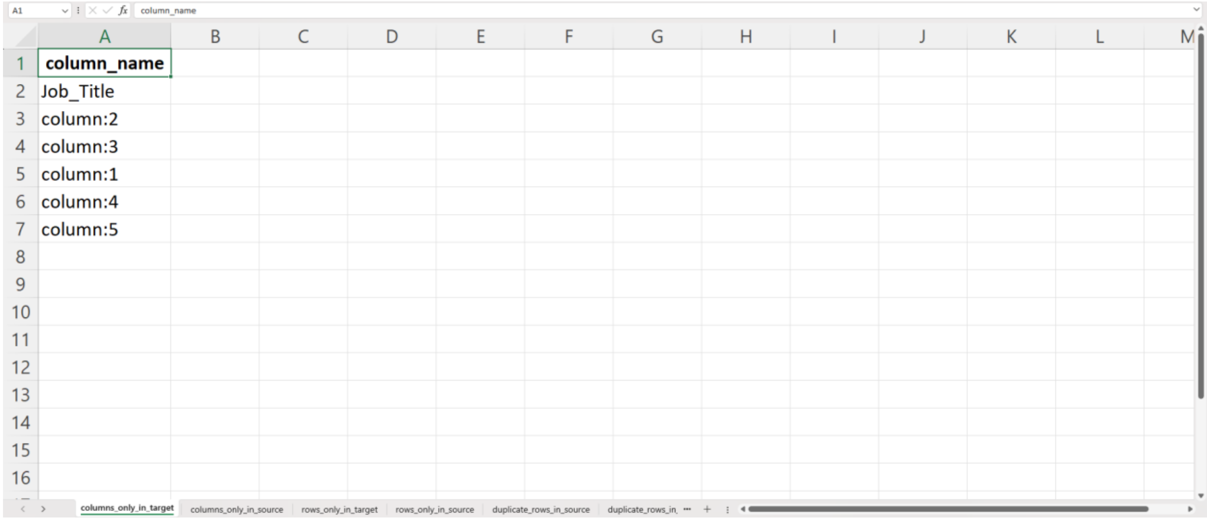

column_name
Job_Title
column:2
column:3
column:1
column:4
column:5

Columns only in source:

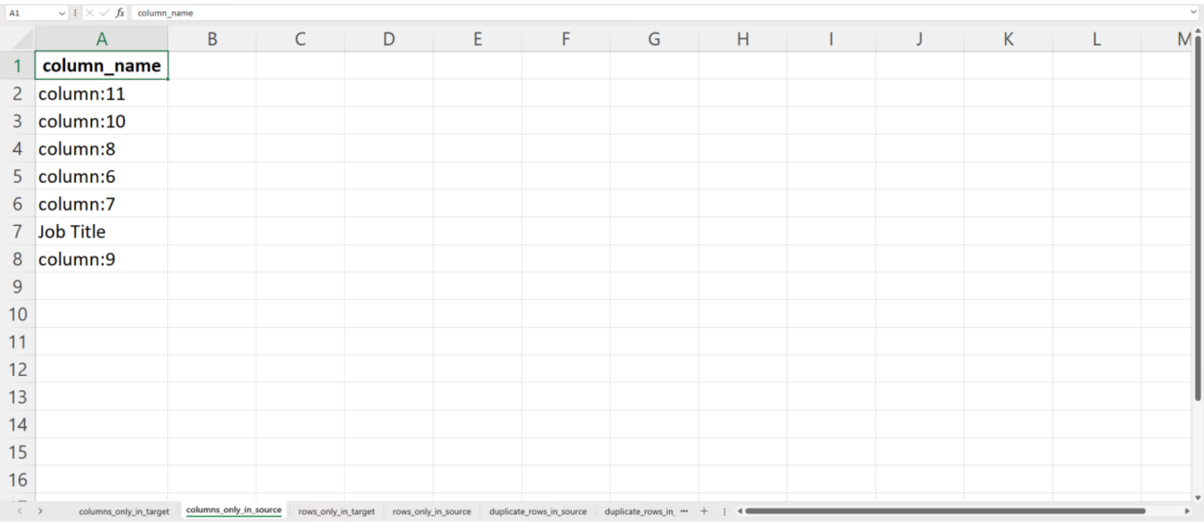

column_name
column:11
column:10
column:8
column:6
column:7
Job Title
column:9

Rows only in target:

Customer no	Name	Email	Telephone	City	Country	Gender	Date Of Birth	Job_Title	column:1	column:
400001	vamsi	something@gmail.com	7895462366	New York	United States	M	09-07-1983	Archivist		

Rows only in source:

Customer no	Name	Email	Telephone	City	Country	Gender	Date Of Birt
1	Danny Hardy	danny.hardy@fake_gmail.com	342-292-3537x5673	New York	United States	M	09-07-1983
1001	Vincent Johnson	vincent.johnson@fake_gmail.com	520.529.7407x05650	New York	United States	M	08-10-2002

Duplicates rows in a source:

Message
There are no records in duplicate_rows_in_source

Duplicates rows in a target:

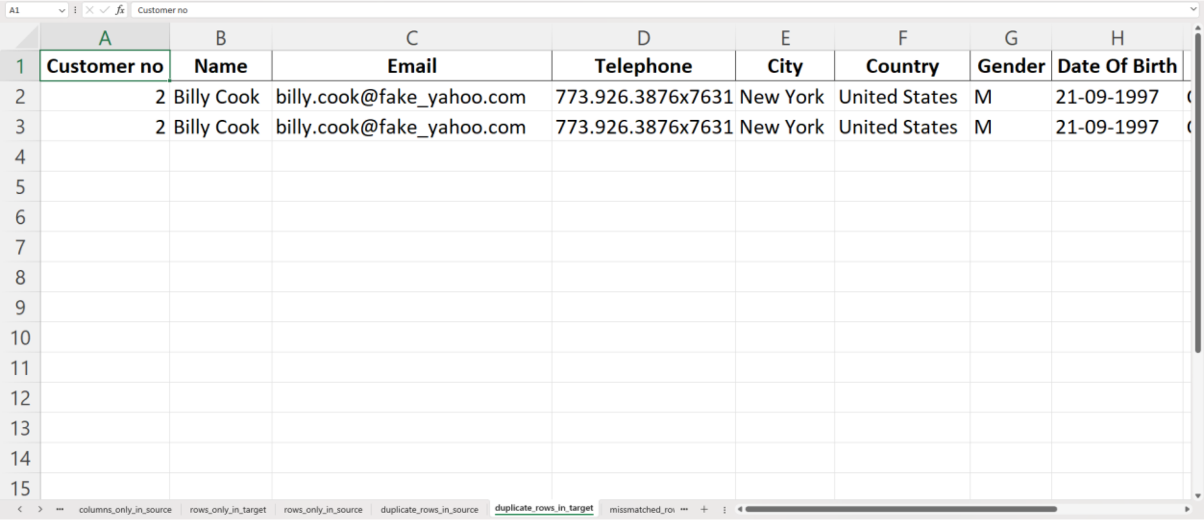

Customer no	Name	Email	Telephone	City	Country	Gender	Date Of Birth
2	Billy Cook	billy.cook@fake_yahoo.com	773.926.3876x7631	New York	United States	M	21-09-1997
2	Billy Cook	billy.cook@fake_yahoo.com	773.926.3876x7631	New York	United States	M	21-09-1997

Miss matched rows from source to target:

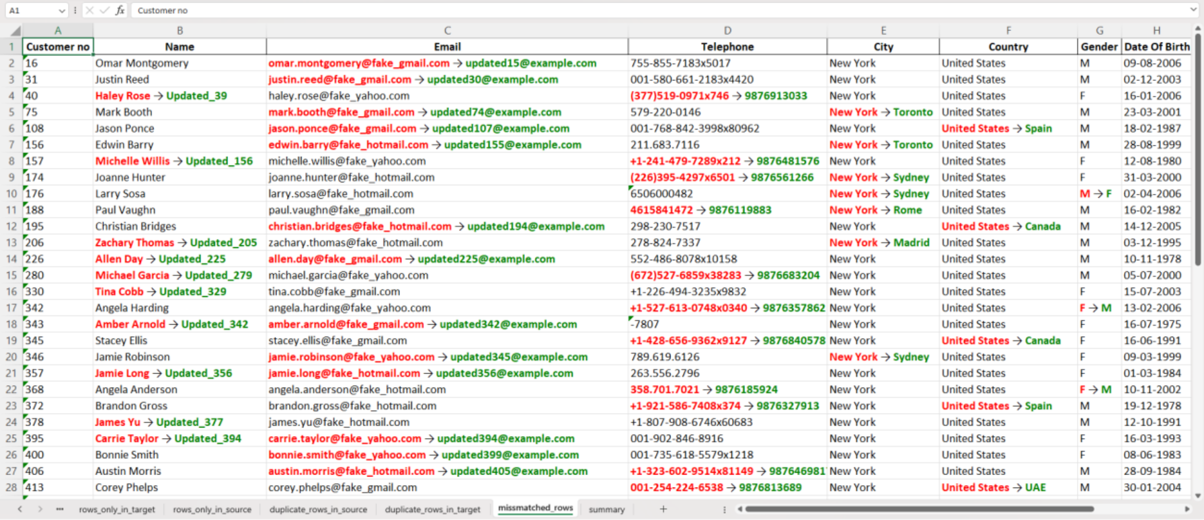

Customer no	Name	Email	Telephone	City	Country	Gender	Date Of Birth
16	Omar Montgomery	omar.montgomery@fake_gmail.com → updated15@example.com	755-855-7183x5017	New York	United States	M	09-08-2006
31	Justin Reed	justin.reed@fake_gmail.com → updated30@example.com	001-580-661-2183x4420	New York	United States	M	02-12-2003
40	Haley Rose → Updated_39	haley.rose@fake_yahoo.com	(377)519-0971x746 → 9876913033	New York	United States	F	16-01-2006
75	Mark Booth	mark.booth@fake_gmail.com → updated74@example.com	579-220-0146	New York → Toronto	United States	M	23-03-2001
108	Jason Ponce	jason.ponce@fake_gmail.com → updated107@example.com	001-768-842-3998x80962	New York	United States → Spain	M	18-02-1987
156	Edwin Barry	edwin.barry@fake_hotmail.com → updated155@example.com	211.683.7116	New York → Toronto	United States	M	28-08-1999
157	Michelle Willis → Updated_156	michelle.willis@fake_yahoo.com	+1-241-479-7289x212 → 9876481576	New York	United States	F	12-08-1980
174	Joanne Hunter	joanne.hunter@fake_hotmail.com	(226)395-4297x6501 → 9876561266	New York → Sydney	United States	F	31-03-2000
176	Larry Sosa	larry.sosa@fake_hotmail.com	6506000482	New York → Sydney	United States	M → F	02-04-2006
188	Paul Vaughn	paul.vaughn@fake_gmail.com	4615841472 → 9876119883	New York → Rome	United States	M	16-02-1982
195	Christian Bridges	christian.bridges@fake_hotmail.com → updated194@example.com	298-230-7517	New York	United States → Canada	M	14-12-2005
206	Zachary Thomas → Updated_205	zachary.thomas@fake_hotmail.com	278-824-7337	New York → Madrid	United States	M	03-12-1995
226	Allen Day → Updated_225	allen.day@fake_gmail.com → updated225@example.com	552-486-8078x10158	New York	United States	M	10-11-1978
280	Michael Garcia → Updated_279	michael.garcia@fake_yahoo.com	(672)527-6859x38283 → 9876683204	New York	United States	M	05-07-2000
330	Tina Cobb → Updated_329	tina.cobb@fake_gmail.com	+1-226-494-3235x9832	New York	United States	F	15-07-2003
342	Angela Harding	angela.harding@fake_yahoo.com	+1-527-613-0748x0340 → 9876357862	New York	United States	F → M	13-02-2006
343	Amber Arnold → Updated_342	amber.arnold@fake_gmail.com → updated342@example.com	-7807	New York	United States	F	16-07-1975
345	Stacey Ellis	stacey.ellis@fake_gmail.com	+1-428-656-9362x9127 → 9876840578	New York	United States → Canada	F	16-06-1991
346	Jamie Robinson	jamie.robinson@fake_yahoo.com → updated345@example.com	789.619.6126	New York → Sydney	United States	F	09-03-1999
357	Jamie Long → Updated_356	jamie.long@fake_hotmail.com → updated356@example.com	263.556.2796	New York	United States	F	01-03-1984
368	Angela Anderson	angela.anderson@fake_hotmail.com	358.701.7021 → 9876185924	New York	United States	F → M	10-11-2002
372	Brandon Gross	brandon.gross@fake_hotmail.com	+1-921-586-7408x374 → 9876327913	New York	United States → Spain	M	19-12-1978
378	James Yu → Updated_377	james.yu@fake_hotmail.com	+1-807-908-6746x60683	New York	United States	M	12-10-1991
395	Carrie Taylor → Updated_394	carrie.taylor@fake_yahoo.com → updated394@example.com	001-902-846-8916	New York	United States	F	16-03-1993
400	Bonnie Smith	bonnie.smith@fake_yahoo.com → updated399@example.com	001-735-618-5579x1218	New York	United States	F	08-06-1983
406	Austin Morris	austin.morris@fake_hotmail.com → updated405@example.com	+1-323-602-9514x81149 → 987646981	New York	United States	M	28-09-1984
413	Corey Phelps	corey.phelps@fake_gmail.com	001-254-224-6538 → 9876813689	New York	United States → UAE	M	30-01-2004

Summary sheet:

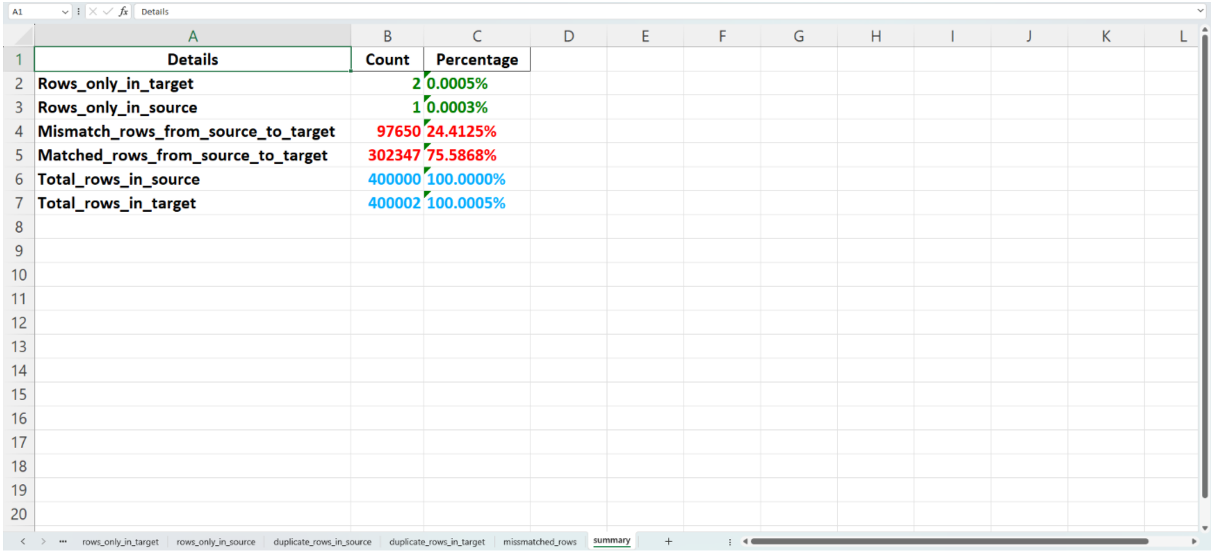

Details	Count	Percentage
Rows_only_in_target	2	0.0005%
Rows_only_in_source	1	0.0003%
Mismatch_rows_from_source_to_target	97650	24.4125%
Matched_rows_from_source_to_target	302347	75.5868%
Total_rows_in_source	400000	100.0000%
Total_rows_in_target	400002	100.0005%

The ETL Validator was able to draw attention to many problems arising from the schema mismatches, missing data, duplicate data, and referential integrity violations in multiple datasets.

2. Text Summary of Results

“ETL Validator handled the task of checking over 1 million records, where clustering of inconsistencies was observed against several validation categories. Schema mismatches occurred in 1.2% of the data; missing values ranged in 1.85% of those 1 million records. The tool effectively caught duplicate records, registering 0.72%, while referential integrity was found missing, by 0.6%. Staying behind in their present functionalities, the ultimate validation took less than 2 minutes," -witnessing high efficiency set, while working with a 1-million-record data set.”

This help enterprises **quickly identify key problem areas in their data pipelines**, ensuring a proactive approach to data quality management.

VI. CONCLUSION

The ETL Validator represents a breakthrough in data integrity management, addressing the critical challenges of modern data transformation processes. By combining advanced algorithmic techniques like MinHash and Locality Sensitive Hashing, the framework provides an unprecedented level of insight into data changes. It transcends traditional validation approaches by offering a nuanced, comprehensive analysis of dataset modifications.

The tool's significance lies not just in detecting differences, but in providing actionable intelligence about data transformations. It empowers data engineers to maintain data quality, identify potential risks, and ensure the reliability of complex ETL pipelines. With its flexible architecture, the validator adapts to diverse data scenarios, making it an indispensable tool for organizations seeking to maintain the highest standards of data integrity and governance.

VII. ACKNOWLEDGMENT

*I wish to **express my heartfelt gratitude to Feuji Software Solutions** for allowing me to take up this internship. It indeed has been a wonderful learning experience, which will help me immensely in my professional and personal development.*

*"I am **extremely thankful to my supervisor and my colleagues for invaluable guidance** and encouragement throughout my time here. His insights and advices have pushed me to the bound of my growth and learning that I thought could never happen. I can't ask for a better mentor.*

I would also like to thank the team that is in Feuji Software Solutions. Thank you for making me so welcome and for answering all my many questions. Your passion and dedication to work are truly inspiring to me."

*I also extend my gratitude to the team at Feuji Software Solutions - **thanks for welcoming me and patiently answering all my queries.** Your love for your job really inspires me.*

This internship has marked my life, and without everyone's help, I could not have done it. Thank you.